

KIKI Agent™

Autonomous LTV Campaign Execution

KIKI Agent™ – Service Level Agreement Framework

Founder: Ariyo Dolapo

Date: 26 July 2026

Status: Confidential — Enterprise Customer SLA

KIKI Agent™ — A STOREGRILL INC LTD Company
Registered in England & Wales

KIKI Agent™ – Service Level Agreement Framework

A STOREGRILL INC LTD Company — Registered in England & Wales

1. Service Commitments

1.1 Uptime Guarantees

Component	Scale Plan	Enterprise Plan
Platform Dashboard & API	99.9% uptime (monthly)	99.95% uptime (monthly)
CAPI Event Ingestion	99.9% delivery rate	99.95% delivery rate
LTV Prediction Engine	99.9% availability	99.95% availability
Fraud Detection Scoring	99.5% availability	99.9% availability
SyncBrain AI Chat	99.5% availability	99.9% availability

1.2 Performance Guarantees

Metric	p50	p99	Measurement Period
LTV Prediction Latency	38ms	95ms	Monthly rolling
CAPI Delivery (Meta)	341ms	500ms	Monthly rolling
CAPI Delivery (All platforms)	400ms	700ms	Monthly rolling
API Response Time	45ms	200ms	Monthly rolling
Dashboard Page Load	1.2s	4s	Monthly rolling

1.3 Financial Remedies

If uptime falls below the committed percentage in any calendar month, the customer receives a Service Credit:

Uptime	Credit (% of monthly fee)
99.5% – 99.9% (Scale) / 99.9% – 99.95% (Enterprise)	0% (within SLA)
99.0% – 99.49% (Scale) / 99.5% – 99.89% (Enterprise)	10% credit
98.0% – 98.99% (Scale) / 99.0% – 99.49% (Enterprise)	25% credit
95.0% – 97.99% (Scale) / 98.0% – 98.99% (Enterprise)	50% credit
Below 95% / Below 98%	100% credit + termination right

Maximum aggregate credit: 100% of one month's fee per rolling 12-month period.

Credit request: Must be submitted within 30 days of the incident month.

2. Support Response Times

2.1 Severity Definitions

Severity	Definition	Examples
Critical (P1)	Platform unavailable. CAPI events not being delivered. LTV predictions failing for all customers.	Dashboard 500 errors, Kafka cluster down, database unreachable
High (P2)	Major feature unavailable. Single customer's events not processing. SyncBrain not responding.	Single platform connector down, fraud engine timeout, SSO failure
Normal (P3)	Non-critical feature issue. Minor UI bug. Documentation error.	Dashboard filter not working, export formatting issue, typo in docs
Low (P4)	Cosmetic issue. Feature request. General enquiry.	Colour mismatch, new platform request, pricing question

2.2 Response & Resolution Targets

Severity	Initial Response	Update Frequency	Target Resolution
P1 — Critical	1 hour (Enterprise) / 2 hours (Scale)	Every 2 hours	4 hours (Enterprise) / 8 hours (Scale)
P2 — High	2 hours (Enterprise) / 4 hours (Scale)	Every 4 hours	8 hours (Enterprise) / 24 hours (Scale)
P3 — Normal	8 business hours	Daily	5 business days
P4 — Low	24 business hours	Weekly	Next release or planned

2.3 Support Channels

Plan	Channels	Hours	Target First Response
Growth	Email only	UK business hours (09:00–18:00, Mon–Fri)	48 hours
Scale	Email + In-app chat	Extended hours (08:00–20:00, Mon–Fri)	8 hours
Enterprise	Email + Chat + Phone + Dedicated Slack	24/7/365	1 hour (P1) / 2 hours (P2)

3. Data Processing & Security

3.1 Data Residency

Data Type	Location	Notes
Customer Conversion Data	Azure UK South (London)	Primary processing region
Enriched Signals (in transit)	Encrypted TLS 1.3	End-to-end to platform endpoints
Dashboard & API Traffic	Azure West Europe (Amsterdam)	CDN-cached where appropriate
Backups	Azure UK South (geo-redundant)	Daily snapshots, 30-day retention
AI Model Training Data	Azure UK South	Anonymised, never leaves UK

3.2 Data Protection

- **Encryption at rest:** AES-256 (Azure Storage Service Encryption)
- **Encryption in transit:** TLS 1.3 minimum (all API endpoints, dashboard, CAPI webhooks)
- **Access control:** Role-based access control (RBAC) — Admin, Operator, Viewer roles
- **Audit logging:** All access to customer data logged and retained for 12 months
- **Data Processing Agreement (DPA):** Available on request, executed before data processing begins
- **Sub-processors:** Azure (infrastructure), Stripe (payments), Groq (AI inference), OpenAI (AI inference) — full list in Privacy Policy

3.3 Backups & Disaster Recovery

Component	Backup Frequency	Recovery Time Objective (RTO)	Recovery Point Objective (RPO)
PostgreSQL Database	Every 6 hours	1 hour	6 hours
File Storage (Azure Files)	Daily	30 minutes	24 hours
Kafka Event Logs	Real-time replication	15 minutes	Near-zero
Application State	Immutable container images	10 minutes	Image-based

DR testing: Conducted quarterly. Customer notification 7 days in advance of scheduled tests.

4. Incident Management

4.1 Incident Classification

Type	Description	Example
Security Incident	Unauthorised access, data breach, credential compromise	AWS key leaked, customer data accessed without authorisation
Availability Incident	Service degradation or outage	Azure region failure, database connection pool exhaustion
Performance Incident	Latency exceeding p99 SLAs	LTV prediction spiking above 200ms
Data Loss Incident	Customer data lost or corrupted	Accidental database truncation, backup failure

4.2 Incident Response Process

1. **Detection** — Automated monitoring (Datadog) or customer report
2. **Triage** — On-call engineer assesses severity (10 minutes)
3. **Containment** — Isolate affected systems, reroute traffic if possible
4. **Resolution** — Fix root cause, verify fix, deploy
5. **Post-Mortem** — Within 48 hours of resolution (Enterprise) / 5 business days (Scale)

4.3 Breach Notification

Under UK GDPR Article 33, KIKI Agent will notify the ICO of any personal data breach within 72 hours of becoming aware. Affected customers will be notified within 24 hours of breach confirmation.

5. Maintenance Windows

Type	Frequency	Notice Period	Downtime Expected	Timing
Routine Platform Updates	Weekly	48 hours	0 minutes (rolling deploy)	Tue/Thu 02:00–04:00 UK
Database Maintenance	Quarterly	7 days	<15 minutes	Weekend 03:00 UK
Infrastructure Upgrade	Bi-annual	14 days	<30 minutes	Weekend 02:00–04:00 UK
Emergency Security Patch	As needed	As soon as possible	Minimised	Immediate

Note: KIKI uses Azure Container Apps with rolling deployments. Most updates involve zero downtime. Any scheduled downtime exceeding 5 minutes will be communicated 48 hours in advance for Scale plans and 7 days in advance for Enterprise.

6. Scope & Exclusions

The SLA does not apply to:

1. Downtime caused by customer-side issues (misconfiguration, expired credentials, network issues at customer site)
 2. Scheduled maintenance (with proper notice)
 3. Force majeure events (Azure region-level outage beyond KIKI's control, natural disasters, war, terrorism)
 4. Third-party platform API downtime (if Meta CAPI or Google Ads API is down, KIKI's delivery SLA may be affected)
 5. Beta or early-access features explicitly marked as "preview" or "experimental"
 6. Free trials, pilots, or non-paying accounts
-

7. Reporting & Verification

- **Monthly SLA reports:** Published within 5 business days of month-end
 - **Real-time status page:** <https://status.kiki.ai>
 - **Verification:** Customers may request third-party monitoring data (Datadog) under NDA
 - **Dispute resolution:** Any SLA credit dispute escalated per Founders' Agreement dispute clause (mediation → arbitration)
-

This SLA framework applies to paid Scale and Enterprise subscriptions. Growth plan customers receive best-effort support with 99.5% target uptime (no financial remedies).

KIKI Agent™ — A STOREGRILL INC LTD Company. Registered in England & Wales.